

ROHAN DUBEY

Lead Data Engineer | Distributed Systems | Real-Time Streaming | Lakehouse & AI/ML Data Platforms

Hyderabad, India | +91-7093123447 | rohankumardubey497@gmail.com | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

SUMMARY

Lead Data Engineer with 8+ years building and operating distributed-data platforms across streaming, batch, lakehouse, and AI/ML workloads. Deep expertise in Flink, Kafka/MSK, Spark, AWS, DynamoDB, S3, Delta Lake, and Iceberg. Led business-critical platforms processing more than 4 billion events daily with 99.99% availability and sub-500 ms latency while delivering approximately \$650K in AWS infrastructure savings.

CORE TECHNICAL SKILLS

Languages: Java, Python, Go, Rust, Scala, SQL, C++, Shell

Streaming & Batch: Apache Flink, Kafka/MSK, Kinesis Data Analytics, Spark, Airflow, event time, watermarks, exactly-once, checkpointing

Cloud & Platform: AWS S3, Glue, EMR, KDA, MSK, Redshift, Athena, Kubernetes, Docker, CI/CD

Storage & Lakehouse: DynamoDB, PostgreSQL, MySQL, Delta Lake, Apache Iceberg, Parquet, ORC, S3 data lake

Engineering Depth: distributed systems, low-latency pipelines, state management, data modeling, schema evolution, observability, JVM/runtime tuning, cost optimization

PROFESSIONAL EXPERIENCE

Flutter Entertainment - Lead Data Engineer, Enterprise Tech

Jun 2026 - Present

Hyderabad, India

- Own enterprise data-platform architecture across four strategic domains: real-time streaming, batch processing, intelligent automation, and AI/ML enablement for global technology functions.
- Define reusable event-driven foundations and cloud architecture standards for reliability, operational efficiency, observability, and business-critical data workflows.
- Led architecture reviews and technical planning across engineering, analytics, and ML teams, translating ambiguous requirements into platform boundaries and influencing a 12-month platform roadmap.

PokerStars - Lead Big Data Engineer, Real-Time Streaming & AI Platforms

Dec 2025 - May 2026

Hyderabad, India

- Led the Streaming AI data-engineering track and mentored five engineers on Flink state management, production debugging, and distributed-systems design.
- Defined a streaming-platform standard adopted by six engineering teams across 20+ production pipelines, establishing reusable patterns for state, checkpointing, exactly-once delivery, recovery, and observability.
- Reduced AWS infrastructure spend by approximately \$650K by eliminating hot-path remote I/O, redesigning keyed-state and storage paths, right-sizing workloads, and tuning checkpoints, sinks, and restarts.

PokerStars - Senior Big Data Engineer

Oct 2023 - Nov 2025

Hyderabad, India

- Engineered and operated real-time pipelines processing more than 4 billion events daily while sustaining 99.99% availability and sub-500 ms processing latency.
- Built reusable platform services for gaming events, customer profiles, personalization, recommendations, operational visibility, analytics, and data science workflows.

Experian PLC - Senior Software Development Engineer

Mar 2023 - Sep 2023

Hyderabad, India

- Delivered backend and Spark processing components for Ascend Intelligence Services, a large-scale SaaS analytics platform spanning loan origination, servicing, and collections.
- Built production components with Java, Spring Boot, Spark SQL, Scala, Akka, Python, Kubernetes, and AWS, translating analytical requirements into maintainable services and pipelines.

Model N - Member of Technical Staff 3

Jul 2021 - Feb 2023

Hyderabad, India

- Designed event-driven data capabilities for an enterprise SaaS revenue-management platform using Spring Boot APIs, Kafka pipelines, Spark processing, Delta Lake, and Dremio analytics.
- Owned backend-to-analytics integration across data movement, processing, lakehouse operations, and governed consumption workflows.

Carelon Global Solutions - Software Developer

Jun 2020 - Jun 2021

Hyderabad, India

- Integrated five healthcare data domains - clinical interventions, membership, claims, providers, and CMS sources - for Medicare/Medicaid analytics and reporting.
- Built aggregation workflows supporting Medicare revenue optimization, administrative-cost reduction, and operational decision-making.

Legato Health Technologies - Associate Software Developer

Jul 2018 - May 2020

Hyderabad, India

- Modernized health-insurance ETL from legacy flat files, Informatica, and Teradata into Parquet, Spark, and AWS data foundations while enabling Tableau reporting.

INDEPENDENT SYSTEMS ENGINEERING

DataWizz ([GitHub](#)): End-to-end lakehouse workspace with three execution engines (DuckDB, PySpark, DataFusion), 10-node visual orchestration, Delta publishing, governance, data quality, notebooks, metrics, alerts, and BI dashboards.

GoXStream ([GitHub](#)): Go streaming engine with REST-defined operator graphs, five operator types, file/database/Kafka connectors, event-time windows, watermarks, late-event handling, and a visual DAG editor.

FlowCore ([GitHub](#)): Rust streaming runtime implementing 10-second tumbling windows, watermarks, late-event detection, periodic checkpoints, event generation, runtime APIs, and a live React dashboard.

EDUCATION

Jawaharlal Nehru Technological University, Hyderabad - B.Tech in Electrical Engineering, GPA 4.0/4.0 (2014 - 2018)